





Qwen 3

(avg. 0.6B, 1.7B, 4B, 8B)



Gemma 3

(avg. 1B, 4B, 12B)



Per-phase quantization

- NVFP4 prefill, BF16 decode
- BF16 prefill, NVFP4 decode
- NVFP4 prefill, NVFP4 decode

Activations quantization

- NVFP4 prefill, NVFP4 decode
- NVFP4 prefill, NVFP4A16 decode
- NVFP4A16 prefill, NVFP4A16 decode