

Pseudo KV-cache compression: accuracy vs rate (sigma = $2^{-\text{bits}} \times \text{RMS}$, groups of 16)

● prefill KV noised ★ decode KV noised no KV noise

